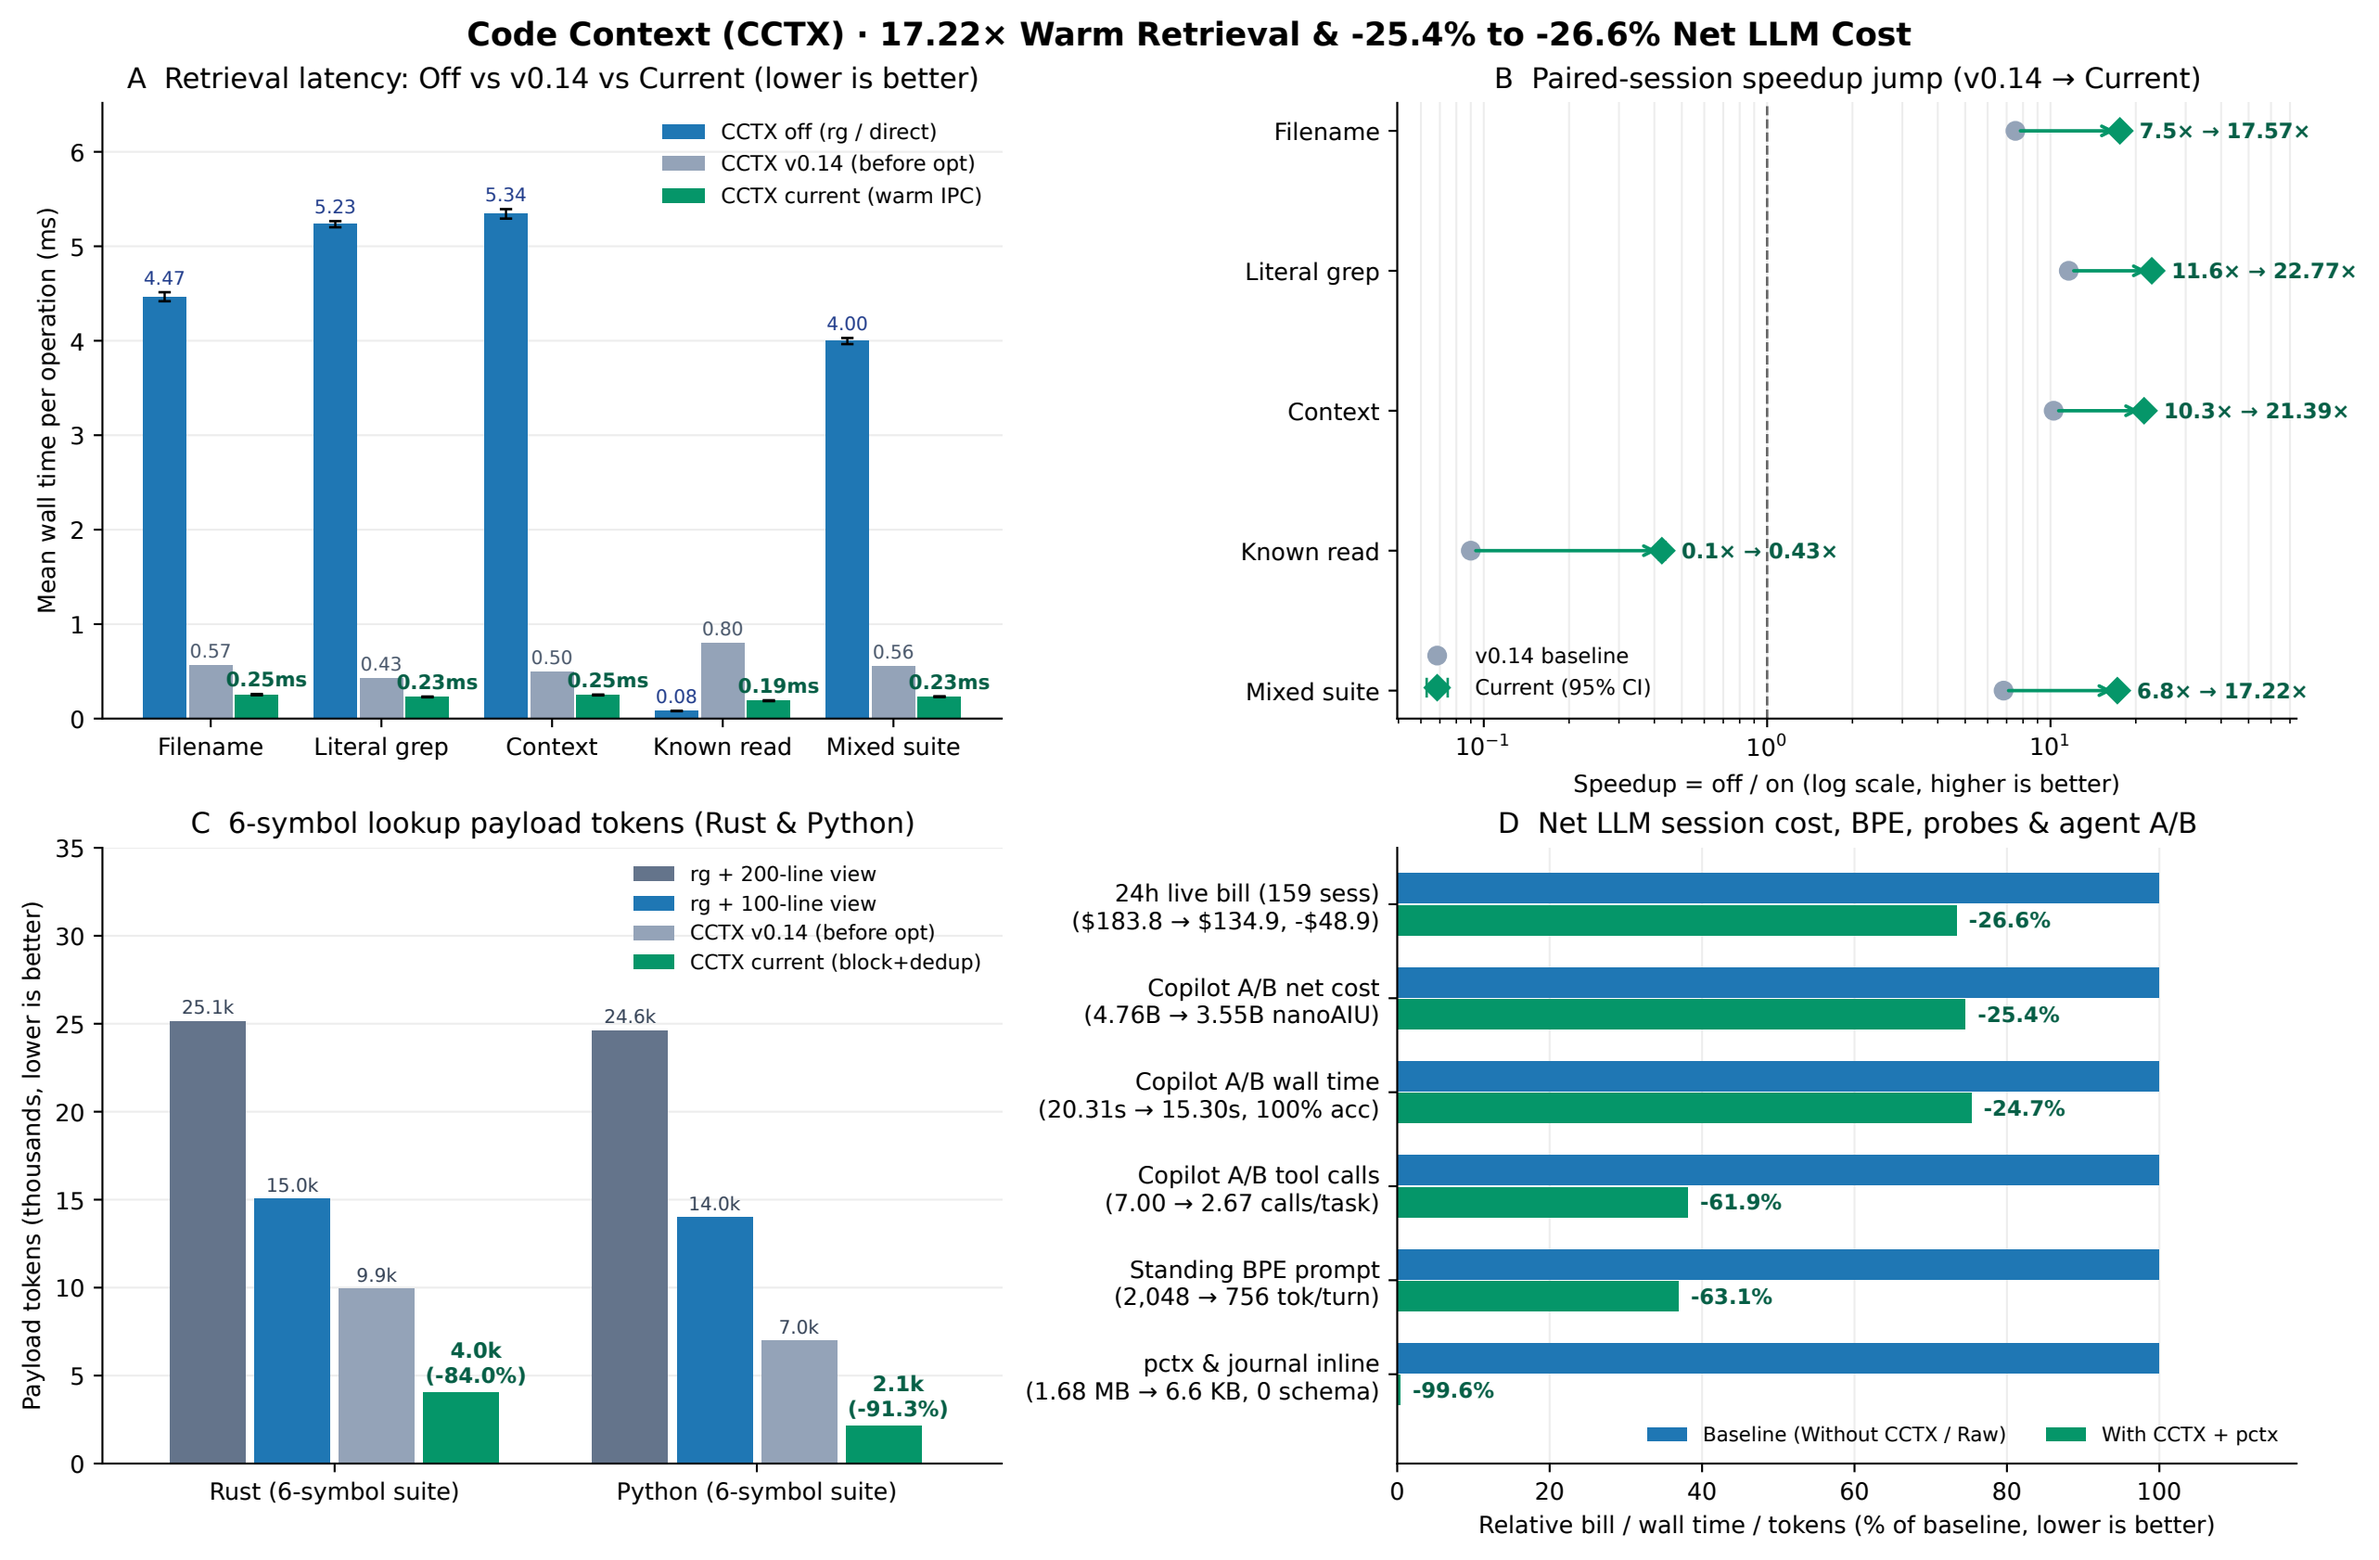
Code Context (CCTX) · 17.22× Warm Retrieval & -25.4% to -26.6% Net LLM Cost
A Retrieval latency: Off vs v0.14 vs Current (lower is better)
CCTX off (rg / direct)
CCTX v0.14 (before opt)
CCTX current (warm IPC)
Mean wall time per operation (ms)
4.47
0.57
0.25ms
5.23
0.43
0.23ms
5.34
0.50
0.25ms
0.08
0.80
0.19ms
4.00
0.56
0.23ms
Filename
Literal grep
Context
Known read
Mixed suite
B Paired-session speedup jump (v0.14 → Current)
Filename
Literal grep
Context
Known read
Mixed suite
7.5× → 17.57×
11.6× → 22.77×
10.3× → 21.39×
0.1× → 0.43×
6.8× → 17.22×
v0.14 baseline
Current (95% CI)
Speedup = off / on (log scale, higher is better)
C 6-symbol lookup payload tokens (Rust & Python)
rg + 200-line view
rg + 100-line view
CCTX v0.14 (before opt)
CCTX current (block+dedup)
Payload tokens (thousands, lower is better)
25.1k
15.0k
9.9k
4.0k
(-84.0%)
24.6k
14.0k
7.0k
2.1k
(-91.3%)
Rust (6-symbol suite)
Python (6-symbol suite)
D Net LLM session cost, BPE, probes & agent A/B
24h live bill (159 sess)
($183.8 → $134.9, -$48.9)
-26.6%
Copilot A/B net cost
(4.76B → 3.55B nanoAIU)
-25.4%
Copilot A/B wall time
(20.31s → 15.30s, 100% acc)
-24.7%
Copilot A/B tool calls
(7.00 → 2.67 calls/task)
-61.9%
Standing BPE prompt
(2,048 → 756 tok/turn)
-63.1%
pctx & journal inline
(1.68 MB → 6.6 KB, 0 schema)
-99.6%
Baseline (Without CCTX / Raw)
With CCTX + pctx
Relative bill / wall time / tokens (% of baseline, lower is better)